

Setup Guide: AI Cement Plant Optimization Platform

Master the deployment of a state-of-the-art, end-to-end AI platform on Google Cloud and Firebase. This definitive guide provides the complete, refined framework for real-time Cement Plant performance monitoring and the automated delivery of intelligent, prescriptive operational intelligence, guaranteeing optimized output across all dynamic operating conditions.

I. Google Cloud Project & Core Setup

First, set up your core GCP project, data storage, and enable necessary services.

1. Create Project and Enable APIs

1. **Create a New GCP Project:**

- Create a new project in the Google Cloud Console. Use a unique ID, such as `genai-cement-optimizer-hack`. We will refer to this as `<PROJECT-ID>`.

2. **Enable Required APIs:**

- Run the following commands in your local terminal (ensure `gcloud` CLI is installed and configured to your `<PROJECT-ID>`):

```
# Core Services
gcloud services enable compute.googleapis.com
gcloud services enable cloudstorage.googleapis.com
gcloud services enable bigquery.googleapis.com

# AI/ML Services
gcloud services enable aiplatform.googleapis.com

# Deployment Services (Cloud Run, Artifact Registry, Firebase)
gcloud services enable cloudbuild.googleapis.com
gcloud services enable run.googleapis.com
gcloud services enable artifactregistry.googleapis.com
gcloud services enable firebase.googleapis.com
```

2. Google Cloud Storage (GCS) - Data and Image Staging

The **raw telemetry data** is staged in GCS, then transformed and made available for Vertex AI via BigQuery. The **raw image data** is also staged in GCS, then transformed and made available for Vertex AI via BigQuery.

1. **Create Data Bucket:**
 - o Navigate to **Cloud Storage**. Click **CREATE BUCKET**.
 - o Name it uniquely (e.g., `cement-optimizer-data-<PROJECT-ID>`).
 - o Select a **Region** (e.g., `us-central1`). **Record this region**, as all subsequent services (BigQuery, Vertex AI, Cloud Run) must use it for optimal performance.
 - o Upload your raw data file (`synthetic_cement_plant_data_v4.csv`) to this bucket.
2. **Create Few-Shot Image Folders (for Gemini):**
 - o In the same bucket (or a new image-specific bucket), create the following folders, which hold the example images required for **Few-Shot Prompting** in the Gemini vision services:
 - `kiln/` (for Kiln image analysis examples)
 - `conveyor/` (for Conveyor Belt image analysis examples)
 - o Upload the relevant example images (e.g., `kiln_operating_normal_1.jpg`, `conveyor_belt_normal.jpg`) into their respective folders.

II. BigQuery - Data Ingestion and Transformation

This prepares the data for the two target KPIs: `clinker_free_lime_pct` and `kiln_specific_thermal_energy`.

1. **Create Dataset:**
 - o Navigate to **BigQuery**. Click your project name, then **CREATE DATASET**.
 - o Name it: `cement_data` and select the **same Region** as your GCS bucket.
2. **Create Table from GCS (Raw Data):**
 - o Create a table named `raw_plant_telemetry` by ingesting your CSV file from GCS, using **Auto detect** for the schema.
3. **Create the Training Data View (Lagged Features):**
 - o This view calculates the necessary features, such as lagged sensor readings.
 - o Click **COMPOSE NEW QUERY** and run your SQL script.
 - o The resulting View, `cement_data.training_data_view`, will be the data source for both AutoML models.

```
CREATE OR REPLACE VIEW `cement_data.training_data_view` AS  
SELECT
```

```

*, -- All original columns
LAG(raw_meal_lsf_ratio, 300) OVER (ORDER BY timestamp) AS raw_meal_lsf_ratio_lag_300
-- Include all other necessary lagged feature calculations here (e.g., ~5 hours of lag)
FROM
`cement_data.raw_plant_telemetry`
WHERE
-- Filter out initial rows where lag values are NULL
raw_meal_lsf_ratio_lag_300 IS NOT NULL;

```

III. Vertex AI - Model Training and Deployment

You will train two models (one for each KPI) and deploy them to their own dedicated endpoints for low-latency predictions.

1. AutoML Training Job (Repeat for 2 Models)

1. **Create Vertex AI Dataset:** (Repeat once for each KPI model)
 - o Navigate to **Vertex AI** \$\rightarrow\$ **Datasets**. Click **CREATE**.
 - o **Name:** E.g., `clinker_free_lime_pct_dataset`.
 - o **Data Source:** Use the BigQuery View URI: `bq://<PROJECT-ID>.cement_data.training_data_view`.
2. **Start Training:** (Repeat once for each KPI model)
 - o Open the dataset and click **TRAIN NEW MODEL**.
 - o **Training Method:** **AutoML**.
 - o **Model Objective:** **Regression**.
 - o **Target column:** Select the target KPI (e.g., `clinker_free_lime_pct` for the first model, then `kiln_specific_thermal_energy` for the second).
 - o **Optimization objective:** Minimize **Root Mean Squared Error (RMSE)**.
 - o Set a reasonable **Training budget** (e.g., 4-6 compute hours).

2. Model Deployment to Endpoint (2 Endpoints)

You must create a separate, dedicated endpoint for each of your **two trained models**.

1. **Deploy Model:** (Repeat for both trained models)
 - o Go to **Vertex AI** \$\rightarrow\$ **Models**. Select a trained model. Go to the **Deploy & test** tab.
 - o Click **Deploy to endpoint**.
 - o **Endpoint name:** Use a clear, descriptive name (e.g., `clinker-free-lime-pct-endpoint`).
 - o **Machine type:** Select a suitable type (e.g., `n1-standard-2`).
 - o **Min/Max compute nodes:** Set both to **1**.

- o Click **Deploy**.
2. **Record Endpoint IDs:**
- o Go to **Vertex AI** \$\rightarrow\$ **Online Prediction** \$\rightarrow\$ **Endpoints**.
 - o **Copy the unique Endpoint ID** for each of your **two deployed models**. These IDs are critical configuration variables for your backend logic.
-

IV. Cloud Run - Backend Services Deployment (5 Services)

You are deploying **five** distinct services in Cloud Run to handle the two-part optimization and two-part image analysis. All services must use the **same region** as your data.

1. Service Account IAM Setup

Your Cloud Run services need permission to interact with Vertex AI, Gemini, and GCS.

1. **Identify Service Account (SA):** The default SA is typically the Compute Engine default: `PROJECT_NUMBER-compute@developer.gserviceaccount.com`.
2. **Grant Roles:** Go to **IAM & Admin** \$\rightarrow\$ **IAM** and grant the following roles to this Service Account:
 - o **Vertex AI User** (`roles/aiplatform.user`) - Allows the services to call the AutoML endpoints and Gemini via Vertex AI.
 - o **Storage Object Viewer** (`roles/storage.objectViewer`) - Allows the Image Analysis services to read the few-shot images from your GCS bucket.

2. Deploy Cloud Run Services

You will run the `gcloud run deploy` command **five times**, once for each distinct microservice.

#	Service Name	Purpose	Notes
1	<code>clinker-kpi-predictor</code>	Predicts <code>clinker_free_lime_pct</code> .	Calls the first AutoML Endpoint.
2	<code>energy-kpi-predictor</code>	Predicts <code>kiln_specific_thermal_energy</code> .	Calls the second AutoML Endpoint.
3	<code>optimization-recommender</code>	Receives 2 KPI predictions, uses Gemini to provide operational recommendations.	Single model handles both KPIs.

#	Service Name	Purpose	Notes
4	kiln-image-analyzer	Uses Gemini (Vision) for Kiln image defect analysis.	Reads from GCS kiln/ folder.
5	conveyor-image-analyzer	Uses Gemini (Vision) for Conveyor Belt image defect analysis.	Reads from GCS conveyor/ folder.

Deployment Command Template (Repeat 5 times, changing the service name):

```
gcloud run deploy <SERVICE-NAME> \
--source . \
--region us-central1 \
--platform managed \
--allow-unauthenticated \
--min-instances 0 \
--max-instances 2
```

3. Record All Service URLs

- After each deployment, **copy the Service URL** from the terminal output. These **five URLs** are the API endpoints that your frontend will call.
-

V. Firebase Hosting - Frontend Deployment

The application frontend is a static HTML/JS file hosted via Firebase.

1. Initialize Firebase Project

1. **Log In and Initialize:**

```
firebase login
firebase init
```

2. **Follow Prompts:**

- Select **Hosting: Configure files for Firebase Hosting**.
- Choose **Use an existing project** and select your GCP project ID.
- **Public directory?** Type `public`.
- **Configure as a single-page app?** Type `N` (No).

2. Prepare and Deploy Frontend

- **Update Cloud Run URLs in `index.html`:**
 - In your local `public/index.html` file, update the JavaScript logic to call the **five Cloud Run Service URLs** you recorded in the previous step. You will have separate POST requests for each service.
- **Deploy the Application:**

```
firebase deploy --only hosting
```

- **Verify and Test:** The output will provide the final **Hosting URL** (e.g., `https://[Your-Project-ID].web.app`).

VI. Technology Stack Summary

This solution leverages a modern, serverless, and decoupled architecture on Google Cloud Platform and Firebase, utilizing specialized services for each phase of the MLOps pipeline and application serving.

GCP/Firebase Service	Category	Role in the Solution
Cloud Run	Serverless Compute	Hosts the five backend microservices : 2 KPI Predictors, 1 Gemini Recommender, 2 Gemini Image Analyzers. Provides automatic scaling from zero.
Vertex AI (AutoML)	Machine Learning	Used for training the two Regression models that predict the critical KPIs.
Vertex AI (Generative AI)	Generative AI	Powers the Optimization Recommender and the Kiln/Conveyor Image Analyzers using multimodal Few-Shot prompting.
Vertex AI (Endpoints)	Machine Learning	Provides managed, low-latency deployment endpoints to serve the predictions from the two trained AutoML models.
BigQuery	Data Warehouse	Stores, manages, and transforms the plant telemetry data, creating the training data view with lagged features.

GCP/Firebase Service	Category	Role in the Solution
Cloud Storage (GCS)	Storage	Serves as the central staging area for raw data, model artifacts, and Few-Shot images.
Firebase Hosting	Frontend Hosting	Provides fast, secure, and globally distributed hosting for the web application frontend.